

Mini Project

Name: Tianying Zhang

CWID: 10466474

Title:

Densely Connected Convolutional Networks

Abstract:

Dense Convolutional Network can be substantially deeper, more accurate, and efficient to train, because they contain shorter connections between layers close to the input and those close to the output. DenseNets, which connects each layer to every other layer in a feed-forward fashion. Whereas traditional convolutional networks with L layers have L connections - one between each layer and its subsequent layer - our network has $L(L+1)/2$ direct connections. For each layer, the feature-maps of all preceding layers are used as inputs, and its own feature-maps are used as inputs into all subsequent layers. DenseNets have several compelling advantages: they alleviate the vanishing-gradient problem, strengthen feature propagation, encourage feature reuse, and substantially reduce the number of parameters.

Github Link: <https://github.com/ZhangxiaoshiHAHAHA/DenseNet>

Introduction:

Based on CIFAR-10 dataset, I build different Densenet structures, such as Densenet121, Densenet169 and Densenet201 on classification task. The accuracy of Densenet121 on test data is 89.95% with 20 epochs.

Method:

A DenseNet is a type of convolutional neural network that utilises dense connections between layers, through Dense Blocks, where we connect all layers (with matching feature-map sizes) directly with each other. To preserve the feed-forward nature, each layer obtains additional inputs from all preceding layers and passes on its own feature-maps to all subsequent layers.

Data:

The CIFAR-10 dataset (Canadian Institute For Advanced Research) is a collection of images that are commonly used to train machine learning and computer vision algorithms. It is one of the most widely used datasets for machine learning research. The CIFAR-10 dataset contains 60,000 32x32 color images in 10 different classes. The 10 different classes represent airplanes, cars, birds, cats, deer, dogs, frogs, horses, ships, and trucks. There are 6,000 images of each class.

In this project, I plan to use 50000 images for training, 10000 images for validation and 10000 images for test.

For all data pre-processing, use `transformes.compose()` combine `ToTensor` and

Normalize functions to deal with it.

Tools & Technologies:

Hardware:

Colab(GPU)

Library:

Pytorch: <https://pytorch.org/>

Numpy: <https://numpy.org/install/>

Matplotlib: <https://matplotlib.org/stable/users/installing/index.html>

Math: <https://pypi.org/project/python-math/>

Experiments:

Training details:

Epoch: 20

Learning rate: e-1

Momentum=0.9,

Weight decay=1e-4

Loss function: Cross Entropy

Optimizer: SGD

Training time: 28mins

Network (Densenet121) :

(3, 32, 32) -> [Conv2d] -> (24, 32, 32) -> [layer1] -> (48, 16, 16) -> [layer2]


->(96, 8, 8) -> [layer3] -> (192, 4, 4) -> [layer4] -> (384, 4, 4) -> [AvgPool]

->(384, 1, 1) -> [Linear] -> (10)

Results:

With parameters set above, the history of training processing: it includes the train_loss, train_acc, test_loss, test_acc and learning rate in each epoch. And it costs about 28 minutes on training.

Epoch [1/ 20]	Train Loss:1.458377	Train Acc:46.82%	Test Loss:1.230144	Test Acc:56.85%	Learning Rate:0.100000
Epoch [2/ 20]	Train Loss:0.912826	Train Acc:67.40%	Test Loss:0.787293	Test Acc:72.08%	Learning Rate:0.100000
Epoch [3/ 20]	Train Loss:0.695529	Train Acc:75.65%	Test Loss:0.842282	Test Acc:71.67%	Learning Rate:0.100000
Epoch [4/ 20]	Train Loss:0.576693	Train Acc:80.01%	Test Loss:0.648330	Test Acc:77.84%	Learning Rate:0.100000
Epoch [5/ 20]	Train Loss:0.505484	Train Acc:82.39%	Test Loss:0.564463	Test Acc:80.67%	Learning Rate:0.100000
Epoch [6/ 20]	Train Loss:0.450149	Train Acc:84.18%	Test Loss:0.510384	Test Acc:82.38%	Learning Rate:0.100000
Epoch [7/ 20]	Train Loss:0.421556	Train Acc:85.43%	Test Loss:0.535721	Test Acc:81.62%	Learning Rate:0.100000
Epoch [8/ 20]	Train Loss:0.387670	Train Acc:86.62%	Test Loss:0.503317	Test Acc:82.89%	Learning Rate:0.100000
Epoch [9/ 20]	Train Loss:0.359841	Train Acc:87.28%	Test Loss:0.455958	Test Acc:84.09%	Learning Rate:0.100000
Epoch [10/ 20]	Train Loss:0.345032	Train Acc:87.96%	Test Loss:0.518581	Test Acc:82.65%	Learning Rate:0.100000
Epoch [11/ 20]	Train Loss:0.239112	Train Acc:92.00%	Test Loss:0.304258	Test Acc:89.67%	Learning Rate:0.010000
Epoch [12/ 20]	Train Loss:0.201749	Train Acc:93.10%	Test Loss:0.310624	Test Acc:89.44%	Learning Rate:0.010000
Epoch [13/ 20]	Train Loss:0.188874	Train Acc:93.62%	Test Loss:0.305213	Test Acc:89.71%	Learning Rate:0.010000
Epoch [14/ 20]	Train Loss:0.179080	Train Acc:93.85%	Test Loss:0.298013	Test Acc:89.82%	Learning Rate:0.010000
Epoch [15/ 20]	Train Loss:0.167836	Train Acc:94.22%	Test Loss:0.294781	Test Acc:90.22%	Learning Rate:0.010000
Epoch [16/ 20]	Train Loss:0.152204	Train Acc:94.86%	Test Loss:0.302152	Test Acc:90.08%	Learning Rate:0.001000
Epoch [17/ 20]	Train Loss:0.148512	Train Acc:95.04%	Test Loss:0.295348	Test Acc:90.13%	Learning Rate:0.001000
Epoch [18/ 20]	Train Loss:0.148494	Train Acc:94.91%	Test Loss:0.284635	Test Acc:90.60%	Learning Rate:0.001000
Epoch [19/ 20]	Train Loss:0.146844	Train Acc:95.02%	Test Loss:0.300247	Test Acc:90.09%	Learning Rate:0.001000
Epoch [20/ 20]	Train Loss:0.143993	Train Acc:95.07%	Test Loss:0.285499	Test Acc:90.55%	Learning Rate:0.001000



Accuracy of the network on the 10000 test images: 89.95 %

Conclusion:

Dense Convolutional Network (DenseNet). It introduces direct connections between any two layers with the same feature-map size. In the experiments, DenseNets tend to yield consistent improvement in accuracy with growing number of parameters without any signs of performance degradation or overfitting. DenseNets may be good feature extractors for various computer vision tasks that build on convolutional features.

Paper Link:

https://openaccess.thecvf.com/content_cvpr_2017/papers/Huang_Densely_Connected_Convolutional_CVPR_2017_paper.pdf